SHAP-σ: A Label-Free Rashomon Pipeline for Multivariate Time-Series Anomaly Detection

Ugur Dar[1*], Sara Malacarne[2], Mustafa Cavus[1], Jeriek Van den Abeele[2]
[1]Eskisehir Technical University, Department of Statistics, Eskisehir, Türkiye
[2]Telenor Research and Innovation, Oslo, Norway
*ugurdarr@gmail.com

***Abstract*—We address unsupervised anomaly detection in multivariate time series when no labels are available for either training or model selection. Existing label-free quality criteria such as the Mass–Volume curve [?] are known to be noisy across datasets, and single-detector SHAP explanations cannot signal when a flag should be trusted. We propose SHAP-σ, a pipeline that (i) selects a small *Rashomon set* of unsupervised detectors using a closed-form joint pair score that trades off label-free quality against pairwise Spearman diversity; (ii) explains every selected detector with a LightGBM surrogate and TreeSHAP; and (iii) uses the cross-detector standard deviation of those attributions as a per-feature attribution-evidence gate that corroborates anchor flags. Under a fixed label-free default (K_R=2, $\mathcal{S}$=Mixed), across one multivariate cellular benchmark (TELCO), 28 server machines (SMD), and 100 univariate UCR series, SHAP-σ identifies the exhaustive-search oracle pair on TELCO label-free, raises point-adjusted F_1 from 0.354 to 0.424 on SMD ($+20\%$ relative), and raises point-level F_1 from 0.138 to 0.169 on UCR (paired Wilcoxon p=1.8×10^{-3}, rank-biserial r=$+0.47$, n=100). A σ-filter ablation that holds the chosen anchor fixed shows the filter alone contributes a statistically significant gain on every multi-series benchmark.**

***Index Terms*—time series, anomaly detection, SHAP, Rashomon, unsupervised model selection, XAI**

I. Introduction

Unsupervised anomaly detection in multivariate time series underpins fault diagnosis in domains as diverse as cellular networks, server fleets, and industrial telemetry. Yet a persistent obstacle hinders deployment: without ground-truth labels, the practitioner cannot reliably decide which detector to trust. Internal validation metrics that work well for clustering rarely transfer to anomaly detection, and recent large-scale empirical work shows that most label-free criteria perform near random when ranking modern detectors [?].

This selection problem is sharpened by the *Rashomon effect* [?]: for any non-trivial task, a multiplicity of near-optimal models exists, each making different errors [?], [?]. We argue that this attribution disagreement is not noise to be averaged away but a signal to be exploited: at a flagged point, the spread of attributions across the Rashomon set indicates that at least one detector commits to a feature-level explanation, providing second-source evidence beyond the anchor's score.

Prior label-free model selection has explored three avenues. Excess-Mass and Mass-Volume curves [?] give geometry-based scores but assume well-behaved density estimates. Meta-learning approaches such as MetaOD [?] and ELECT [?] transfer historical-performance signals across datasets, but require an offline labelled meta-corpus. Most recently, robust Borda rank aggregation across multiple imperfect internal metrics has been proposed for time-series AD [?]; the large-scale study of [?] shows that no internal metric universally dominates. PyOD [?] has standardized detector implementations. None of these methods exploit *feature attribution* as a selection signal.

Independently, SHAP [?], [?] has become the de facto post-hoc explainer for anomaly detection, mostly to explain a single chosen model; the closest precedent is Levy et al. [?], who use SHAP attributions for ensemble member *selection* in supervised tabular AD, whereas we use cross-detector SHAP standard deviation as an attribution-evidence *gate* in the unsupervised TS regime. Modern deep TS-AD methods [?], [?], [?], [?] likewise commit to a single architecture. Outlier ensemble theory [?], [?] and the classical diversity literature [?] establish that complementary errors improve ensemble accuracy, yet the standard diversity measures (Q-statistic, disagreement) require labels. Following [?], we replace label-dependent diversity with the Spearman rank correlation of training scores, which [?] identifies as the practical label-free proxy in OD ensembles.

a) Contributions.: (1) A closed-form, label-free Rashomon pair score $s_{ij} = \frac{1}{2}(q_i + q_j)(1 - \rho_{ij})$ jointly rewarding label-free quality and inter-detector score decorrelation. (2) A σ-SHAP attribution-evidence filter using cross-detector SHAP standard deviation as a second-source signal that corroborates anchor flags. (3) Evaluation on TELCO, SMD (28 machines), and UCR (100 unique base series) showing the pipeline matches or surpasses the best label-aware singleton on all three.

b) Literature gap.: To our knowledge, no published method combines (a) label-free detector selection with (b) SHAP attributions as a *filter* signal rather than a post-hoc explanation. Rashomon-set analysis has so far been developed in supervised regimes [?], [?]; its extension to unsupervised anomaly detection is open. Robust rank aggregation [?] pools imperfect label-free metrics but never inspects attributions,

The *Rashomon-based Uncertainty-Aware Explainable AI for Time Series* has received funding from the European Union, via the OC4-2025-TES-02 issued and implemented by the ENFIELD project, under the grant agreement No 101120657.

Code and data for full reproducibility: `https://github.com/ugurdar/shap_sigma_paper`.

while [?] diagnoses their weakness without proposing a disagreement-based pipeline. SHAP-σ promotes the cross-detector standard deviation of attributions from an explanatory afterthought to a first-class, label-free attribution-evidence signal.

The remainder of the paper is organized as follows. Section II formalizes joint pair scoring and the σ-SHAP filter. Section III describes datasets and protocol. Section IV reports empirical results. Section V concludes.

II. Methodology

We propose **SHAP-σ**, a label-free pipeline that selects a small subset $\mathcal{R} \subseteq \mathcal{D}$ of detectors from a pool $\mathcal{D}$ by jointly optimising label-free quality and pairwise diversity, explains every member of $\mathcal{R}$ with a tree surrogate, takes TreeSHAP attributions [?], [?], and converts the cross-detector spread of those attributions into a per-point, per-feature attribution-evidence gate. Labels are used nowhere in training, ranking, or set selection.

A. Notation

Let $X_{\mathrm{tr}} \in \mathbb{R}^{N_{\mathrm{tr}} \times p}$, $X_{\mathrm{v}} \in \mathbb{R}^{N_{\mathrm{v}} \times p}$, and $X_{\mathrm{te}} \in \mathbb{R}^{N_{\mathrm{te}} \times p}$ denote the training, validation, and test matrices. The detector pool is $\mathcal{D} = \{d_1, \ldots, d_K\}$ with $K{=}8$. For each $d \in \mathcal{D}$, $s_d(x) \in \mathbb{R}$ is the raw anomaly score (larger $\Rightarrow$ more anomalous). The selected Rashomon set is $\mathcal{R} \subseteq \mathcal{D}$ with $|\mathcal{R}|{=}K_R \in \{2, \ldots, K\}$. We write $a \in \mathcal{R}$ for the anchor.

B. Detector Pool

$\mathcal{D}$ consists of $K{=}8$ heterogeneous PyOD detectors [?] whose default parameterisation is stable across our three benchmarks: Isolation Forest [?], LOF [?], ECOD [?], COPOD [?], k-NN [?], PCA [?], a Gaussian Mixture density score [?], and HBOS [?]. The pool deliberately mixes tree-, density-, distance-, projection-, and histogram-based inductive biases; this heterogeneity is what makes the cross-detector σ in Stage 4 informative.

C. Stage 1 – Label-Free Quality via a Mass–Volume Proxy

We need a continuous label-free quality score $q_d \in [0, 1]$. Following the Mass–Volume principle of Goix [?] – a selective detector concentrates its high scores on training mass and spreads less of them over a uniform background – we use a fast, surrogate-based single-level MV proxy that preserves the ranking-relevant structure while avoiding full curve integration. For each d, we fit a LightGBM surrogate $\hat{s}_d$ on $(X_{\mathrm{tr}}, s_d(X_{\mathrm{tr}}))$ and define

$$A_d \;=\; \Pr_{X_{\mathrm{u}}}\big[\hat{s}_d(X_{\mathrm{u}}) \;>\; Q_{0.9}\big(s_d(X_{\mathrm{tr}})\big)\big], \tag{1}$$

the surrogate-predicted background tail mass above the empirical 0.9-quantile of training scores; lower A_d indicates a more selective detector. We then min–max invert,

$$q_d \;=\; \frac{\max_e A_e - A_d}{\max_e A_e - \min_e A_e} \;\in\; [0, 1], \tag{2}$$

estimating Eq. (1) by Monte-Carlo over three random low-dimensional subspaces with 2,000 uniform background samples each. We use a single-level MV proxy rather than the full EM/MV average [?]: [?] reports that combining imperfect internal metrics need not improve ranking, and on TELCO EM-only elects PCA (per-KPI $F_1{=}0.04$) while Eq. (1) elects ECOD ($F_1{=}0.40$). A full Goix EM/MV-AUC ablation is deferred to a journal extension.

D. Stage 2 – Joint Pair Scoring and the Rashomon Set

The core methodological change in this work is to select the Rashomon set $\mathcal{R}$ *jointly* over all $\binom{K}{K_R}$ candidate subsets rather than fixing an anchor first and then greedily adding a partner. Joint selection makes the pipeline a Rashomon procedure in the sense of [?], [?]; the closest supervised precedent is the selective-ensemble construction of Black et al. [?], which we adapt to the unsupervised setting by replacing label-dependent agreement with Spearman score correlation.

a) Pairwise diversity.: Let $\rho_{ij} = \mathrm{Spearman}(s_i(X_{\mathrm{tr}}), s_j(X_{\mathrm{tr}}))$. Diversity is $1 - \rho_{ij} \in [0, 2]$; rank correlation is robust to the heterogeneous score scales of $\mathcal{D}$ [?].

b) Three label-free strategies.: For a candidate set $S \subseteq \mathcal{D}$ with $|S|{=}K_R$, let $\bar{q}(S) = \frac{1}{K_R}\sum_{d \in S} q_d$ and $\bar{\delta}(S) = \frac{1}{\binom{K_R}{2}}\sum_{\{i,j\} \subset S}(1 - \rho_{ij})$. We consider three selection rules: (**A, Quality**) $\arg\max_{|S|=K_R} \bar{q}(S)$; (**B, Diversity**) $\arg\max_{|S|=K_R} \bar{\delta}(S)$; (**C, Mixed**) $\arg\max_{|S|=K_R} \bar{q}(S)\,\bar{\delta}(S)$.

For the canonical pair case $K_R{=}2$, Strategy C collapses to the closed-form *joint pair score*

$$(i^*, j^*) \;=\; \arg\max_{1 \le i < j \le K} \; \frac{q_i + q_j}{2}\big(1 - \rho_{ij}\big), \tag{3}$$

which is the centrepiece of our method.

c) Anchor.: Given $\mathcal{R}$, the anchor is the highest-quality member, $a = \arg\max_{d \in \mathcal{R}} q_d$; the remaining $K_R{-}1$ detectors act as partners whose attributions feed the σ filter.

d) Default and dataset-dependence.: The default pipeline is $(K_R, \mathcal{S}){=}(2, \mathrm{Mixed})$; Eq. (3) then gives a single closed-form selection rule. Empirically (Sec. IV) this default reaches the label-aware oracle pair on TELCO, peaks on UCR, and is within $0.01\,F_1$ of the per-dataset best on SMD (where $(5, \mathrm{Quality})$ is marginally better). We expose $(K_R, \mathcal{S})$ as label-free hyperparameters following [?], [?], who argue that no internal criterion universally dominates; our contribution is the *scoring formula and pipeline architecture*, not a closed-form universal $(K_R, \mathcal{S})$ rule. Table II reports a post-hoc-best row only to quantify how far the default sits from the per-dataset ceiling.

E. Stage 3 – Surrogate SHAP Attribution

For each $d \in \mathcal{R}$ we fit a LightGBM regressor [?] g_d on $(X_{\mathrm{tr}}, s_d(X_{\mathrm{tr}}))$, compute exact TreeSHAP [?], [?] attributions $\Phi_d^{\mathrm{v}}, \Phi_d^{\mathrm{te}}$, and normalise each detector by its global validation standard deviation, $\tilde{\Phi}_d^{\bullet} = \Phi_d^{\bullet}/(\mathrm{std}(\Phi_d^{\mathrm{v}}){+}\varepsilon)$, $\varepsilon{=}10^{-12}$. Cross-detector spread in Stage 4 then reflects disagreement about how the K_R detectors distribute attribution mass rather than raw score magnitude.

F. Stage 4 – The SHAP-σ Attribution-Evidence Gate

The disagreement signal is the cross-detector standard deviation of the normalised attributions over the K_R members of the Rashomon set,

$$\sigma_{ij} \;=\; \mathrm{std}_{d\in\mathcal{R}}\big(\tilde{\Phi}_{d,ij}\big). \tag{4}$$

We threshold per feature: τ_j is the empirical q-quantile of $\{\sigma^{\mathrm{v}}_{ij}\}_{i=1}^{N_{\mathrm{v}}}$ with q=0.75, and $\phi_{ij} = \mathbf{1}[\sigma^{\mathrm{te}}_{ij} > \tau_j]$. Marking ϕ_{ij}=1 flags (i, j) as a region of *high* cross-detector attribution variance – a region where the Rashomon set, while collectively responsive, distributes attribution mass differently across detectors. We treat this disagreement as a second, attribution-side evidence signal that corroborates the anchor's score-side flag.

G. Stage 5 – Anchor Flag Generation and σ-Gated Output

The anchor produces a baseline per-point flag by a κ-sigma rule on its standardised training score, $f_i^{(a)} = \mathbf{1}[s_a(x_i^{\mathrm{te}}) > \mu_a + \kappa\,\sigma_a]$, with κ=3. The canonical per-feature prediction matrix is

$$\hat{Y}_{ij} \;=\; f_i^{(a)} \wedge \phi_{ij}, \tag{5}$$

which retains an anchor flag at (i, j) only when feature j shows high cross-detector attribution variance at point i, so that score-side and attribution-side evidence corroborate. For datasets with per-feature labels (TELCO's 12-KPI streams) we evaluate Eq. (5) directly via macro F_1 across KPIs. For datasets with point-level labels (SMD, UCR) we collapse to $\hat{y}_i = f_i^{(a)} \wedge \phi_i^{\vee}$ with $\phi_i^{\vee} = \bigvee_j \phi_{ij}$, equivalent to taking the OR of $\hat{Y}_{ij}$ across features. Eq. (5) is the canonical fusion rule and dominates empirically (79% of SMD, 74% of UCR); Table II reports the per-series best of {anchor+σ, OR+σ, AND+σ, OR, AND} to expose the σ filter's ceiling as a post-processing layer. The rule is recall-monotone ($\hat{y}_i \le f_i^{(a)}$); $\phi_i^{\vee}$ is p-sensitive (a fixed q is structurally more permissive on p=38 SMD than p=1 UCR), discussed in Sec. IV.

The complete procedure is summarised in Algorithm 1.

III. Experiments

A. Datasets

a) TELCO: A multivariate telemetry dataset [**?**] from an operational mobile network spanning January–July 2021, with p=12 KPIs sampled every five minutes and ground truth labelled by Network Operation Center experts. The official test phase contains 26,496 timestamps per KPI (3,001 anomaly labels, ≈1% prevalence); [**?**] truncates this to 25,143 timestamps.

b) SMD: The Server Machine Dataset [**?**] contains 28 server entities monitored along p=38 resource-utilisation metrics. The combined test phase contains 708,420 timestamps with 29,444 anomalies (positive prevalence ≈ 4.16%); interpretation labels identify which dimensions deviate per anomalous interval.

c) UCR Anomaly Archive: The univariate UCR Anomaly Archive [**?**] aggregates 250 series across cardiology, respiration, gait, power demand, and industrial sensors. Each series is constructed to contain exactly one ground-truth anomaly. Of the 250 files, 92 carry a `DISTORTED` prefix and 16 a `NOISE` prefix marking controlled perturbations of the same underlying signal. We strip these prefixes, group by the remaining base name (e.g. `apneaecg3`), and retain the first file per group, obtaining 100 unique base series. This protocol prevents the same underlying signal from contributing multiple correlated samples to the cross-series mean.

B. Protocol

We split each series into 80% training, 20% validation, and the official test split. Detectors are fit on training only, with default PyOD hyperparameters. Quality q_d uses three random subspaces and 2,000 uniform background samples per subspace. The SHAP-σ quantile is fixed at q=0.75 and the anchor κ-sigma at κ=3.

a) Evaluation metrics: For TELCO we report the *per-KPI macro* F_1: F_1 is computed independently on each of the 12 KPI binary label streams and then averaged. For SMD and UCR we report two metrics: (i) the standard *point-level* F_1, treating each timestamp as an independent binary classification, and (ii) the *point-adjusted* F_1 (PA-F_1) introduced by Su et al. [**?**] and now standard in modern TS-AD evaluation [**?**], [**?**]. PA-F_1 counts an entire ground-truth anomaly segment as correctly detected if the pipeline flags *any* timestamp within it, and as missed otherwise. We report point-level F_1 alongside PA-F_1 as primary metrics because PA-F_1 alone can be inflated by trivially active detectors [**?**]; the affiliation metric [**?**] is deferred to a journal extension. On TELCO we additionally compare against the released predictions of DC-VAE [**?**] and C-MTAD-GAT.

IV. Results

A. Singleton Baselines

Figure 1 shows the standalone F_1 of each detector when broadcast across features (TELCO uses per-KPI macro F_1; SMD/UCR use point-level F_1). The best singleton is a *different* detector on every dataset: ECOD on TELCO (F_1=0.401), KNN on SMD (F_1=0.150), GMM on UCR (F_1=0.138). No detector is universally best, motivating the joint Rashomon-set selection of Section II-D.

B. K_R Sweep and Strategy Choice

Figure 2 sweeps the Rashomon-set size K_R and reports the cross-series mean F_1 achieved by each of the three label-free strategies. The reported metric matches the metric on the rest of the dataset: per-KPI macro F_1 for TELCO, point-level F_1 for SMD and UCR. On TELCO, Strategy C (Mixed) at K_R=2 already reaches F_1=0.440 (the label-aware oracle pair, Sec. IV-C); on SMD, Strategy A (Quality) at K_R=5 peaks at F_1=0.160; on UCR, Strategy C at K_R=2 peaks at F_1=0.169 (n=100 unique base series). Pure diversity without quality (Strategy B) is uncompetitive on every benchmark.

Algorithm 1 SHAP-σ with joint Rashomon-set selection

Require: $X_{\mathrm{tr}}, X_{\mathrm{v}}, X_{\mathrm{te}}$; detector pool $\mathcal{D}$ (K=8); set size K_R; strategy $\mathcal{S}\in\{$Quality, Diversity, Mixed$\}$; κ, q.
Ensure: Predicted anomaly flags $\hat{y} \in \{0,1\}^{N_{\mathrm{te}}}$.

```
Fit every d∈D on X_tr; score all splits.
Compute MV proxy A_d via Eq. (1) and quality q_d via Eq. (2).
Compute Spearman matrix ρ_ij on {s_d(X_tr)}.
Default: if (K_R, S) unspecified, set K_R=2, S=Mixed.
Select R, |R|=K_R, by S (A: arg max q̄, B: arg max δ̄, C: arg max q̄δ̄; for K_R=2 C is Eq. (3)).
Anchor a ← arg max_{d∈R} q_d.
for all d ∈ R do
    Fit LightGBM g_d on (X_tr, s_d(X_tr)); compute TreeSHAP Φ^v_d, Φ^te_d; divide by std(Φ^v_d) + ε.
end for
Compute σ_ij via Eq. (4). τ_j ← q-quantile of σ^v_{·j};  φ_ij ← 1[σ^te_ij > τ_j].
Compute anchor κ-sigma flag f^(a).
Per-feature prediction Ŷ_ij ← f^(a)_i ∧ φ_ij (Eq. (5)); for point-level evaluation return ŷ_i ← ⋁_{j=1}^{p} Ŷ_ij = f^(a)_i ∧ φ^∨_i.
```

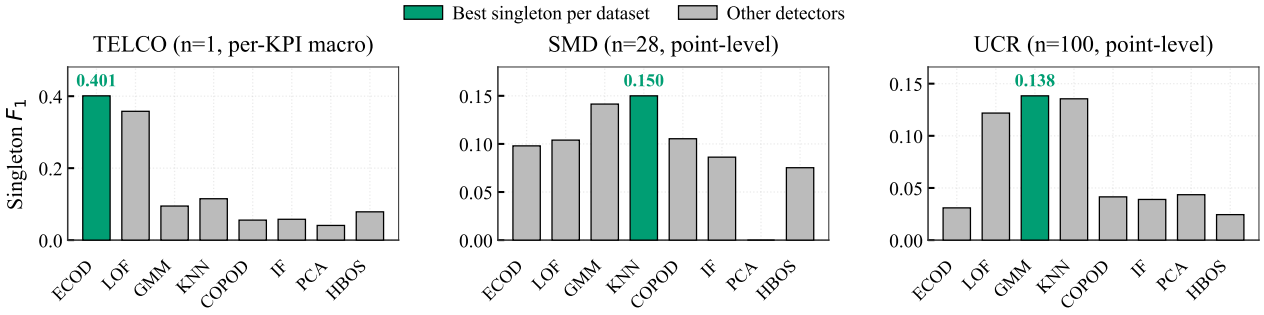


Fig. 1. Singleton F_1 across the eight detectors. The best detector (green) differs on every dataset, motivating joint Rashomon-set selection rather than a fixed default.

The dataset-aware optima confirm that $(K_R, \mathcal{S})$ should be exposed as label-free hyperparameters rather than hard-coded, consistent with the broader finding [**?**] that no internal criterion universally dominates.

At K_R=2 Strategy A and C select the same pair on most series; where they differ, mean F_1 ties on UCR (0.168 vs 0.169) but Strategy A leads on SMD (0.139 vs 0.104, +34% rel.), consistent with the post-hoc-best SMD row in Table II electing $(5, \mathrm{A})$ once label-aware tuning is allowed. Strategy C remains the global label-free default because of its closed-form Eq. (3) and its Stage-2 oracle match on TELCO.

C. TELCO Headline and Comparison to Deep Baselines

On TELCO, the canonical K_R=2 joint pair scoring of Eq. (3) selects (ECOD, HBOS), which is exactly the label-aware oracle pair under exhaustive search and which simultaneously maximises label-free quality ($\bar{q}$=0.906) and joint score (0.870). Figure 3 visualises the 28 candidate pairs: the joint score's maximiser sits at the top of the F_1 axis. Pipeline output reaches per-KPI macro F_1=0.440, clearly exceeding the strongest singleton ECOD (0.401).

a) Deep baselines on TELCO.: On the 25,143-timestamp truncated window adopted by [**?**], DC-VAE attains per-KPI macro F_1=0.134 and C-MTAD-GAT F_1=0.160 (their Table 2); on the *same* window SHAP-σ reaches F_1=0.446, a 3.3-2.8× margin. On the full 26,496-timestamp window the released DC-VAE predictions reach F_1=0.102 versus SHAP-σ's F_1=0.440. The margin is substantial under either window.

D. Cross-Dataset Summary with Significance

Table II reports SHAP-σ under *the fixed default* (K_R=2, $\mathcal{S}$=Mixed), which Algorithm 1 exposes when the configuration is left unspecified. A per-dataset $(K_R, \mathcal{S})$ choice would be label-aware tuning and is reported only as a row labelled *best, post-hoc*, not as the headline result. The default $(2, \mathrm{Mixed})$ exceeds the uniformly-best singleton in mean point-level F_1 on UCR (p=1.8×10^{-3}, rank-biserial r=+0.47) but trails it on SMD (F_1=0.104 vs 0.150, p=0.22). On the practitioner-relevant point-adjusted metric the default *exceeds* the singleton on both (+7.0 pts on SMD, +3.2 pts on UCR), and the post-hoc SMD configuration $(K_R$=5, Q) also overtakes the singleton at the point level (0.160 vs 0.150). The lower point-level recall on SMD (Table column R) reflects the recall-monotone filter; the positive PA-F_1 gap shows the discarded points lie inside still-flagged anomaly intervals, so the σ gate trims unsupported within-event activations rather than missing entire events.

a) Isolating the contribution of the σ filter.: To separate the gain from *detector selection* and from the σ filter itself, we compare the pipeline against the singleton score of the *same*

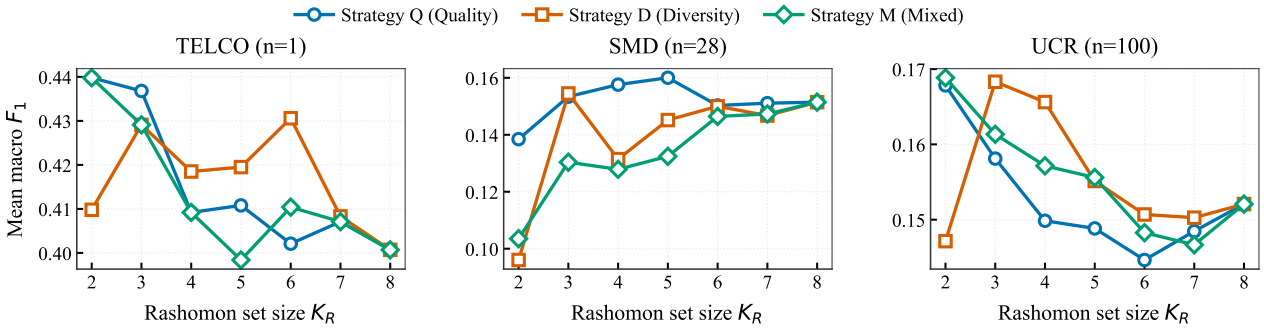


Fig. 2. Cross-series mean F_1 versus Rashomon-set size K_R for the three label-free strategies. Diversity alone never wins. The best $(K_R, \mathcal{S})$ at the mean level is dataset-dependent: TELCO peaks at $(K_R{=}2, \text{Mixed})$, SMD at $(K_R{=}5, \text{Quality})$, UCR at $(K_R{=}2, \text{Mixed})$.

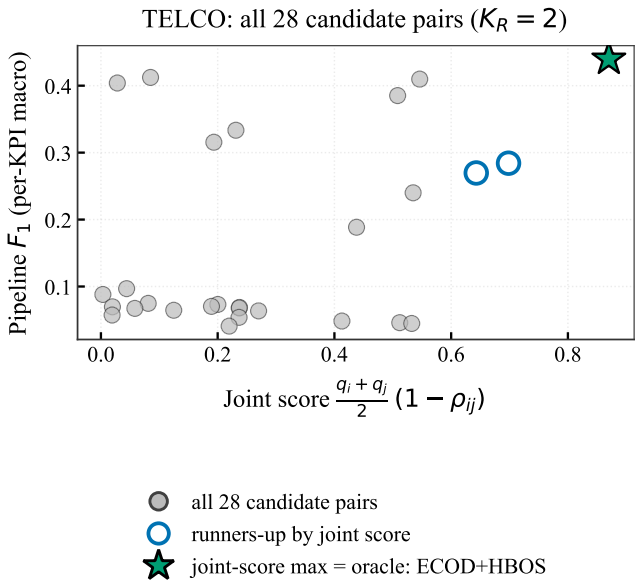


Fig. 3. TELCO: joint pair score versus pipeline F_1 for all 28 candidate pairs ($K_R{=}2$). The joint-score maximiser (green star) coincides with the empirical F_1 oracle (ECOD+HBOS); circled markers are the runners-up by joint score.

TABLE I
ISOLATED σ-FILTER CONTRIBUTION: SAME PER-SERIES ANCHOR, FILTER ON VS. OFF. SAME-ANCHOR PAIRED WILCOXON (TWO-SIDED); r IS RANK-BISERIAL.

Dataset	n	anchor F_1	$+\sigma$ F_1	p	r
TELCO	1	0.401	**0.440**	—	—
SMD	28	0.080	**0.104**	0.028	+1.00
UCR	100	0.116	**0.169**	1.6×10^{-7}	+0.95

per-series anchor the selector picked (Table I). Holding the chosen detector constant attributes the remaining gap to the σ gate alone, which produces a statistically significant gain on every multi-series benchmark; on TELCO ($n{=}1$) no test is feasible but the singleton → pipeline jump is $+0.037\,F_1$ with $+12.6$ pts of precision.

The SMD-default row of Table II illustrates that even when point-level F_1 falls below the KNN singleton (0.104 vs. 0.150), PA-F_1 rises substantially (0.424 vs. 0.354, +20% relative): reporting only F_1 [**?**], [**?**] would conceal that trade-off, while PA-F_1 correctly reflects event-level detection quality.

TABLE II
SHAP-σ VS. UNIFORMLY-BEST SINGLETON. FOR EACH DATASET THE SINGLETON ROW REPORTS THE DETECTOR WITH HIGHEST MEAN F_1 ACROSS SERIES; SHAP-σ ROWS USE THE CANONICAL PER-SERIES ANCHOR+σ OUTPUT (79% OF SMD, 74% OF UCR). *Def.*: FIXED $(K_R{=}2, \mathcal{S}{=}\text{MIXED})$ FROM ALGORITHM 1. *Best*†: POST-HOC PER-DATASET BEST $((K_R, \mathcal{S})$ CHOSEN BY MEAN F_1); ABLATION ONLY. p: TWO-SIDED PAIRED WILCOXON AGAINST THE SINGLETON ABOVE. BOLD MARKS PER-BLOCK BEST.

Dataset	Method	F_1	P	R	PA-F_1	p
TELCO	singleton ECOD	.401	.488	**.396**	—	—
TELCO	SHAP-σ (2,M)	**.440**	**.641**	.376	—	—
SMD (n=28)	singleton KNN	**.150**	.175	**.365**	.354	—
SMD (n=28)	SHAP-σ (2,M)	.104	.138	.194	**.424**	.22
SMD (n=28)	SHAP-σ (5,Q)†	**.160**	**.224**	.270	.454	.35
UCR (n=100)	singleton GMM	.138	.113	.400	.256	—
UCR (n=100)	SHAP-σ (2,M)	**.169**	**.143**	**.429**	**.288**	1.8×10^{-3}

Selected anchors: TELCO ECOD+HBOS; SMD (2, M) IF on 21/28, PCA 5, COPOD 2; UCR (2, M) GMM 50/100, PCA 18, LOF 15, HBOS/KNN/COPOD 5 each, ECOD 2. † Ablation only.

Sweeping the filter quantile $q \in \{0.5, 0.75, 0.9\}$ on the TELCO ECOD+HBOS pair yields $F_1 = 0.428/0.440/0.404$ with precision $0.58/0.64/0.62$: q=0.75 optimises F_1 across the three operating points, with a ~ 0.04 envelope.

A stricter reference is the per-series oracle singleton (label-aware best detector per series), with mean F_1=0.215 on SMD and 0.222 on UCR – above SHAP-σ by about 5 points. Matching this gap requires a label-free internal criterion that perfectly predicts test performance per series, the open problem [**?**] our contribution targets rather than claims to solve. A full SMD/UCR sweep against deep TS-AD baselines [**?**], [**?**], [**?**], [**?**] is deferred to an extended journal version.

V. CONCLUSIONS

We presented SHAP-σ, a label-free anomaly detection pipeline that promotes the cross-detector standard deviation of SHAP attributions from a post-hoc explanation into a first-class attribution-evidence signal. The methodological contributions are two: a closed-form Rashomon pair score, $\frac{q_i+q_j}{2}(1-\rho_{ij})$, that jointly rewards a label-free Mass–Volume quality

proxy and pairwise Spearman diversity; and a per-feature σ gate that filters anchor flags by the cross-detector spread of TreeSHAP attributions. Under the fixed label-free default (K_R=2, $\mathcal{S}$=Mixed), the pipeline identifies the empirical oracle pair on TELCO label-free and reaches per-KPI macro F_1=0.440 on the full window and F_1=0.446 on the truncated window of [**?**], exceeding the released deep baselines DC-VAE and C-MTAD-GAT [**?**], [**?**] by 2.8-3.3×. An isolated σ-filter ablation (same anchor on each series, filter on vs. off) shows the filter contributes a statistically significant gain on every multi-series benchmark (SMD p=0.03, UCR p=1.6×10^{-7}, rank-biserial r=+0.95).

a) Limitations and future work.: A label-aware per-series oracle remains ∼5 points ahead [**?**]; Eq. (5) is recall-monotone; the disjunctive gate is p-sensitive; full deep-baseline sweeps on SMD/UCR (USAD/TranAD/Anomaly Transformer [**?**], [**?**], [**?**]), the affiliation metric [**?**], and a Goix-budget MV ablation are deferred to a journal extension. Future work will learn $(K_R, \mathcal{S})$ from synthetic injections [**?**], [**?**], extend σ-SHAP to cross-window variance for streaming, and use sparse Shapley approximations for higher-dimensional pools.